

National University of Singapore
School of Computing

CS2105

Assignment 2

Semester 1 AY23/24

Introduction and Submission Deadline

In this assignment, you will implement a reliable one-way chat program¹ that sends messages over an unreliable UDP channel that may either corrupt or drop packets randomly (but will always deliver packets in order).

The deadline for submission is **30 October 2023 (Monday), 11:59 pm** sharp. Do not leave your submission to the last minute to avoid unforeseeable complications.

2 points penalty will be imposed on late submissions (i.e. submissions or re-submissions made after the deadline). **No submission will be accepted after 6 November 2023 (Monday), 11:59 pm and 0 points will be awarded.**

Group Work

All the work in this assignment should be done individually. However, if you find the assignment too difficult,

- you are allowed to form a group with another student
- **maximum two students per group.**
- Group submission is subject to **a 2 points penalty** for each student.

Under no circumstances should you solve it in a group and then submit it as an individual solution. This is considered plagiarism. **There will be no acceptance for excuses such as forgetting to declare as group submission.** Please refer to the "Special Instructions for Group Submission" and "Plagiarism Warning" on page 3 for more details.

Grading

We will test and grade your program on the stu server. Please make sure that your programs run correctly on stu. Moreover, you can use libraries installed in public folders of stu (e.g. /usr/lib) only.

We accept submission of Python 3 (**3.10**), Java, or C/C++ program, and we recommend that you use **Python 3** for your assignments. Programming languages other than Python 3, Java, and C/C++ are not allowed. For Python 3, we use the python3 program installed in folder /usr/bin on stu for grading. If you use Java or C/C++, we will compile and run your program for grading using the default compilers on stu (openjdk

¹We implement a one-way chat to keep the assignment simple.

17 installed in /usr/local/jdk/bin, or gcc 11.4.0 installed in /usr/bin). For C/C++ users, we will link your program with zlib during compilation. The grading script infers your programming language from the file extension name (.py, .java, .c). Therefore, please ensure your files have the correct extension names. For a Java program, the class name should be consistent with the source file name, and please implement the static `main()` method so that the program can be executed as a standalone process after compilation.

We will grade your program based on its correctness only. A grading script will be used to test your program, and no manual grading will be provided.

Program Submission

For individual submission, please submit a zip file that contains source programs and submit it to the Assignment 2 assignment on Canvas.

- Please name your zip file as **<Matric number>.zip**, where **<Matric number>** refers to your matric number that starts with the letter A. Note that your student number's first and last letters should be capital letters. One example file name would be **A0112345X.zip**.
- Your zip file should **only** contain **two source programs**, **Alice.py** and **Bob.py**. If you use Java, C, or C++, replace ".py" with ".java", ".c", or ".cpp", respectively. Do not put your programs under any subfolders (or package paths).
- For resubmissions, use the same file name, and allow Canvas to append numbers to the file.

Special Instructions for Group Submission

- Designate one person to submit the zip file for group work instead of two different persons submitting the same file twice.
- The file name should contain the student numbers of two members and be named as **<Matric number 1>-<Matric number 2>.zip**. An example would be **A0165432X-A0123456Y.zip**.
- **Do not change the designated submitter!** If the group needs to upload a new version, it should be done by the same designated submitter as well.

We will **deduct 1 mark** for every type of failure to follow instructions (e.g. wrong program name, wrong zip file name, folder structure).

Grading Rubric

We will grade your programs based on their correctness only. A grading script will be used to test your programs, and **no manual grading** will be provided. The grading script infers your programming language from the file extension name (.py, .java, .c). Therefore, please ensure your files have the correct extension names. For a Java program,

the class name should be consistent with the source file name. Please implement the **public static main()** method so that the program can be executed as a standalone process after compilation.

1. **[1 point]** Programs are free from syntax errors and can be successfully executed (or compiled for Java/C/C++). Program execution follows the specified commands strictly. Source files are correctly named, and there are **no additional subfolders inside the zip file**.
2. **[1 point]** Programs can successfully send chat messages from Alice to Bob in a perfectly reliable channel (i.e. no error at all).
3. **[2 points]** Programs can successfully send messages from Alice to Bob in the presence of **packet corruptions**.
4. **[2 points]** Programs can successfully send messages from Alice to Bob in the presence of **packet losses**.
5. **[1 point]** Programs can successfully send messages from Alice to Bob in the presence of **both packet corruptions and packet losses**.

NOTE:

- The chat messages received by Bob must be **identical** to the ones sent by Alice to claim a successful transmission.
- The total size of input messages to Alice is guaranteed to be no larger than **5,000 bytes** (including newline characters).
- During grading, Alice and Bob processes will be forced to terminate **10 seconds** after Alice is started. The grading script will then compare the output of Bob against the input of Alice. Please ensure your programs are efficient enough to meet the above criteria.

CAUTION: The grading script **grades your program according to what Bob prints out on the screen**. Thus, make sure that you **remove all debug messages** before you submit your solution. **No points will be awarded in each test case if your output does not conform to the expected result.**

Question and Answer

If you have any doubts about this assignment, please post your questions on LumiNUS forum before consulting the teaching team. **We will not debug programs for you.** However, we may help to clarify misconceptions or give necessary directions if required.

Plagiarism Warning

You are free to discuss this assignment with your friends. However, ultimately, you should write your own code. We employ a zero-tolerance policy against plagiarism. If a suspicious case is found, we will award zero points and may take further disciplinary action.

A Word of Advice

This assignment can be time-consuming. We suggest that you start writing programs **early, incrementally and modularly**. For example, deal with error-free transmission first, then data packet corruption, ACK packet corruption, etc. Test your programs frequently and make backup copies before every significant change. Frequent backups will allow you to submit at least a partially working solution that is worth some marks.

Do not post your solution in any public domain on the Internet or share it with friends, even after this semester.

Overall Architecture

There are three programs involved in this assignment, Alice, UnreliNET and Bob (see Figure 1 below).

- Alice will send chat messages to Bob over UDP (and Bob may provide feedback as necessary).
- The UnreliNET program is used to simulate the unreliable transmission channel (in both directions) that randomly corrupt or drop packets. You can assume that UnreliNET always delivers packets in order.
- Instead of sending packets directly to Bob, Alice will send all packets to UnreliNET. UnreliNET may introduce bit errors or lose packets randomly. It then forwards packets (if not lost) to Bob.
- When receiving feedback packets from Bob, UnreliNET may also corrupt them or lose them with a certain probability before relaying them to Alice.



Figure 1: UnreliNET Simulates Unreliable Network

The UnreliNET program is complete and provided in the assignment. Your task is to develop the Alice and Bob programs so that Bob will receive chat messages successfully in the presence of packet corruption and packet loss.

UnreliNET Program

The UnreliNET program simulates an unreliable channel that may corrupt or lose packets with tunable probability. **This program is complete and provided as a Java class file.** There's no necessity to view and understand the source code of UnreliNET to complete this assignment.

To run UnreliNET on sunfire, type the following command:

```
java UnreliNET <P_DATA_CORRUPT> <P_DATA_LOSS>  
<P_ACK_CORRUPT> <P_ACK_LOSS> <unreliNetPort> <rcvPort>
```

For example:

```
java UnreliNET 0.3 0.2 0.1 0.05 9000 9001
```

The above command makes UnreliNET listen on port 9000 of localhost and forwards all received data packets to Bob running on the same host with port 9001, with 30% chance of packet corruption and 20% chance of packet loss. The UnreliNET program also forwards ACK/NAK packets back to Alice, with a 10% packet corruption rate and a 5% packet loss rate.

Note that UnreliNET only accepts packets with a **maximum payload** of **64 bytes (inclusive of user-defined header/trailer fields)**. It will drop any packet beyond this size limit and show an error message on the screen.

Packet Error Rate

The UnreliNET program randomly corrupts or drops data packets and ACK/NAK packets according to the specified parameters `P_DATA_CORRUPT`, `P_DATA_LOSS`, `P_ACK_CORRUPT`, and `P_ACK_LOSS`. Please choose these values in the range `[0, 0.3]` during testing (setting a too large corruption/loss rate may result in a prolonged transmission). The grading script also chooses parameters in the range `[0, 0.3]`.

If you have trouble getting your code to work, set the parameters to 0 first for debugging purposes.

Alice Program

- Alice will read chat messages from **standard input** and send them to UnreliNET (which will then forward to Bob as appropriate).
- The chat messages contain **ASCII characters only**, and there is no empty message (i.e. blank line).
- There's **no delay** regarding the input to Alice.

To run Alice on `stu`, type the following command:

```
python3 Alice.py <unreliNetPort>
```

`<unreliNetPort>` is the port number UnreliNET is listening to. For example:

```
python3 Alice.py 9000
```

Bob Program

- Bob receives chat messages from Alice (via UnreliNET) and prints them to **standard output**.
- It may also send feedback packets to Alice (also via UnreliNET) as you deem appropriate.

The command to run the Python version of Bob is:

```
python3 Bob.py <rcvPort>
```

For example:

```
python3 Bob.py 9001
```

With the above command, Bob listens to port 9001 of `localhost` and prints all received messages to the standard output.

Some tips are given below:

1. Bob should only print messages that are correctly received from Alice, no more, no less.
2. Make sure you do not print anything irrelevant to standard output. In particular, remember to **remove all debug messages before submission**.
3. You may use standard error (stderr in Python) for debugging/logging purposes.

Running All Three Programs

For this assignment, you will always run `UnreliNET`, `Alice` and `Bob` programs on the **same host**.

1. You should launch `Bob` first.
2. Followed by `UnreliNET` in the second window.
3. Finally, launch `Alice` in a third window to start data transmission.

The `UnreliNET` program simulates an unreliable network and runs infinitely. Once launched, you may reuse it in consecutive tests if you do not want to change its parameters. To manually terminate it, press `<Ctrl> + c`. Do terminate any unwanted running instance before starting the next run or exiting `stu`.

- The `Alice` and `Bob` programs **should not** communicate directly – all traffic has to go through the `UnreliNET` program.
- `Alice` should terminate once all input is read from the user and adequately forwarded (i.e. the input stream is closed and `Bob` successfully receives everything in the input stream).
- There is no need for `Bob` to detect the end of transmission and terminate. If you need to manually terminate it, press `<Ctrl> + c`.

Testing Your Programs

To test your program, please use your SoC UNIX ID and password to log on to `stu` as instructed in Assignment 0.

If you test your server manually, please select a port number greater than 1024 because the OS usually restricts the usage of ports lower than 1024. If you get a `BindException: Address already in use` (or similar errors for other languages), please try a different port number.

For the convenience of testing, you can use the file redirection feature to let `Alice` read from a file (rather than from keyboard input).

For example, you can run `Alice` as follows:

```
python3 Alice.py <unreliNetPort> <input.txt>
```

You can replace `input.txt` with another file name if you wish.

In the same way, you can let `Bob` print to a file (rather than print on screen - just for the

convenience of testing):

```
python3 Bob.py <rcvPort> > output.txt
```

Remember to flush the output buffer so that data will be written to output.txt.

Finally, you can compare input.txt to output.txt by issuing the following command:

```
cmp input.txt output.txt
```

This line compares two files byte by byte and prints the differences if found. If the two files are binary equivalent, nothing will be printed on the screen. This method is useful to detect unexpected outputs such as hidden characters.

A sample input.txt is provided for your testing.

We also released a comprehensive set of grading scripts for you. There are **hidden test** cases during grading. Passing all the test cases **does not guarantee** that you will get full marks. We will detect fraudulent cases such as **hard-coding** answers, or Alice communicates with Bob in **any means other than via UnreliNET**.

To use the grading script, please upload your programs and the test folder given in the package to stu. Make sure that your programs and the test folder are **in the same directory**. Then, you can run the following command to test your programs:

```
bash test/RunTest.sh
```

By default, the script runs through all test cases. You can also choose to run a certain test case by specifying the case number in the command:

```
bash test/RunTest.sh 2
```

To stop a test, press <Ctrl> + c. If pressing the key combination once does not work, hold the keys until the script exits.

Note: If you receive "Failed" feedback indicating wrong output especially smaller output size, it can be due to timeout. In this case, you might also see some program exceptions followed by "KeyboardInterrupt" before the "Failed" line. This might be because your programs are forced to be terminated by the grading program due to timeout.

Self-defined Header/Trailer Fields at Application Layer

UDP transmission is unreliable. To detect packet corruption or packet loss, you need to implement reliability checking and recovery mechanisms yourself. The following header/trailer fields are recommended though you may choose a different design:

1. Sequence number
2. Checksum

You are reminded again that each packet Alice or Bob sends should contain **at most 64 bytes of payload data (inclusive of user-defined header/trailer fields)**, or UnreliNET will drop it.

Computing Checksum

To detect bit errors, Alice should compute the checksum for **every outgoing packet**. Please refer to Assignment 0 exercise 1 on how to compute the checksum.

Timer and Timeout Value

Alice may have to maintain a timer for unacknowledged packets. You should set a socket timeout value of **50ms**. Socket timeout can be set using Python function `socket.settimeout()`. Other languages provide similar features, which can be found online.